

Keras/Tensorflow Modelling Tutorial: Iris Dataset

This tutorial will demonstrate using Keras for deep learning on the Iris dataset. We will also go through the following:

- Creating data and target matrices from imported data
- One hot encoding of non-numerical class labels
- Building a convolutional neural network (CNN) model
- Accuracy after training and testing the model
- Depiction of confusion matrix

This code is adapted from "Multi-Class Classification Tutorial with the Keras Deep Learning Library" from Machine Learning Mastery and "Deep Learning Iris Dataset Keras" from Kaggle. The links are included below:

- <https://machinelearningmastery.com/multi-class-classification-tutorial-keras-deep-learning-library/>
- <https://www.kaggle.com/akashsri99/deep-learning-iris-dataset-keras>

Importing Libraries

In [1]:

```
import warnings; warnings.simplefilter('ignore')

import os
os.environ['KERAS_BACKEND'] = 'tensorflow'

from keras.backend import set_image_dim_ordering
set_image_dim_ordering('tf')
```

Using TensorFlow backend.

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from seaborn import heatmap
%matplotlib inline
```

In [3]:

```
from keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD, Adam
```

In [4]:

```
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from keras.utils import np_utils
```

Creating data (X) and target (Y) matrices

In [5]:

```
df = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data", header=None)
df.head()
```

Out[5]:

	0	1	2	3	4
--	---	---	---	---	---


```
[0., 0., 1.],
[0., 0., 1.],
[0., 0., 1.]], dtype=float32)
```

Building Learning Model

For this example, we will build a one layer network with 4 inputs, 8 hidden nodes, and 3 outputs for demonstration.

In [17]:

```
from sklearn.model_selection import train_test_split
```

In [18]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=8)
```

In [38]:

```
model = Sequential()
```

In [39]:

```
model.add(Dense(8, input_shape=(4,), activation='relu'))
model.add(Dense(3, activation='softmax'))
```

In [40]:

```
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 8)	40
dense_17 (Dense)	(None, 3)	27

Total params: 67
Trainable params: 67
Non-trainable params: 0

In [41]:

```
model.compile(Adam(lr=0.04), 'categorical_crossentropy', metrics=['accuracy'])
```

In [42]:

```
history = model.fit(X_train, y_train, epochs=100)
```

```
WARNING: Logging before flag parsing goes to stderr.
W0801 03:27:00.197965 140735736558528 deprecation.py:323] From
/Users/xiaosg/anaconda3/lib/python3.6/site-packages/tensorflow/python/ops/math_grad.py:1250:
add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

```
Train on 105 samples
Epoch 1/100
105/105 [=====] - 0s 2ms/sample - loss: 1.5491 - accuracy: 0.2667
Epoch 2/100
105/105 [=====] - 0s 92us/sample - loss: 1.0354 - accuracy: 0.4571
Epoch 3/100
105/105 [=====] - 0s 121us/sample - loss: 0.8507 - accuracy: 0.6571
Epoch 4/100
105/105 [=====] - 0s 152us/sample - loss: 0.5320 - accuracy: 0.7429
```

```
Epoch 5/100
105/105 [=====] - 0s 162us/sample - loss: 0.4578 - accuracy: 0.7619
Epoch 6/100
105/105 [=====] - 0s 131us/sample - loss: 0.4315 - accuracy: 0.7714
Epoch 7/100
105/105 [=====] - 0s 98us/sample - loss: 0.3555 - accuracy: 0.8857
Epoch 8/100
105/105 [=====] - 0s 134us/sample - loss: 0.2837 - accuracy: 0.9429
Epoch 9/100
105/105 [=====] - 0s 127us/sample - loss: 0.2634 - accuracy: 0.9619
Epoch 10/100
105/105 [=====] - 0s 157us/sample - loss: 0.2367 - accuracy: 0.9429
Epoch 11/100
105/105 [=====] - 0s 136us/sample - loss: 0.2044 - accuracy: 0.9714
Epoch 12/100
105/105 [=====] - 0s 145us/sample - loss: 0.1960 - accuracy: 0.9619
Epoch 13/100
105/105 [=====] - 0s 166us/sample - loss: 0.1825 - accuracy: 0.9524
Epoch 14/100
105/105 [=====] - 0s 158us/sample - loss: 0.1604 - accuracy: 0.9810
Epoch 15/100
105/105 [=====] - 0s 151us/sample - loss: 0.1487 - accuracy: 0.9619
Epoch 16/100
105/105 [=====] - 0s 162us/sample - loss: 0.1550 - accuracy: 0.9429
Epoch 17/100
105/105 [=====] - 0s 158us/sample - loss: 0.1531 - accuracy: 0.9429
Epoch 18/100
105/105 [=====] - 0s 130us/sample - loss: 0.1598 - accuracy: 0.9333
Epoch 19/100
105/105 [=====] - 0s 183us/sample - loss: 0.1417 - accuracy: 0.9619
Epoch 20/100
105/105 [=====] - 0s 146us/sample - loss: 0.1225 - accuracy: 0.9810
Epoch 21/100
105/105 [=====] - 0s 194us/sample - loss: 0.1347 - accuracy: 0.9524
Epoch 22/100
105/105 [=====] - 0s 160us/sample - loss: 0.0998 - accuracy: 0.9810
Epoch 23/100
105/105 [=====] - ETA: 0s - loss: 0.1129 - accuracy: 0.96 - 0s
160us/sample - loss: 0.1023 - accuracy: 0.9810
Epoch 24/100
105/105 [=====] - 0s 150us/sample - loss: 0.0983 - accuracy: 0.9810
Epoch 25/100
105/105 [=====] - 0s 217us/sample - loss: 0.0960 - accuracy: 0.9810
Epoch 26/100
105/105 [=====] - 0s 173us/sample - loss: 0.0904 - accuracy: 0.9714
Epoch 27/100
105/105 [=====] - 0s 140us/sample - loss: 0.1065 - accuracy: 0.9619
Epoch 28/100
105/105 [=====] - 0s 117us/sample - loss: 0.0864 - accuracy: 0.9714
Epoch 29/100
105/105 [=====] - 0s 175us/sample - loss: 0.0913 - accuracy: 0.9810
Epoch 30/100
105/105 [=====] - 0s 159us/sample - loss: 0.0894 - accuracy: 0.9810
Epoch 31/100
105/105 [=====] - 0s 224us/sample - loss: 0.0932 - accuracy: 0.9714
Epoch 32/100
105/105 [=====] - 0s 167us/sample - loss: 0.0965 - accuracy: 0.9619
Epoch 33/100
105/105 [=====] - 0s 177us/sample - loss: 0.0818 - accuracy: 0.9810
Epoch 34/100
105/105 [=====] - 0s 202us/sample - loss: 0.0948 - accuracy: 0.9619
Epoch 35/100
105/105 [=====] - 0s 192us/sample - loss: 0.0801 - accuracy: 0.9714
Epoch 36/100
105/105 [=====] - 0s 200us/sample - loss: 0.0755 - accuracy: 0.9810
Epoch 37/100
105/105 [=====] - 0s 173us/sample - loss: 0.0718 - accuracy: 0.9810
Epoch 38/100
105/105 [=====] - 0s 187us/sample - loss: 0.0768 - accuracy: 0.9714
Epoch 39/100
105/105 [=====] - 0s 181us/sample - loss: 0.0722 - accuracy: 0.9810
Epoch 40/100
105/105 [=====] - 0s 284us/sample - loss: 0.0745 - accuracy: 0.9810
Epoch 41/100
105/105 [=====] - 0s 172us/sample - loss: 0.0784 - accuracy: 0.9810
Epoch 42/100
105/105 [=====] - 0s 148us/sample - loss: 0.0863 - accuracy: 0.9714
```

```
Epoch 43/100
105/105 [=====] - 0s 151us/sample - loss: 0.0794 - accuracy: 0.9810
Epoch 44/100
105/105 [=====] - 0s 184us/sample - loss: 0.0696 - accuracy: 0.9810
Epoch 45/100
105/105 [=====] - ETA: 0s - loss: 0.1051 - accuracy: 0.96 - 0s
164us/sample - loss: 0.0805 - accuracy: 0.9714
Epoch 46/100
105/105 [=====] - 0s 172us/sample - loss: 0.0695 - accuracy: 0.9810
Epoch 47/100
105/105 [=====] - 0s 190us/sample - loss: 0.0719 - accuracy: 0.9714
Epoch 48/100
105/105 [=====] - 0s 147us/sample - loss: 0.0701 - accuracy: 0.9810
Epoch 49/100
105/105 [=====] - 0s 139us/sample - loss: 0.0635 - accuracy: 0.9810
Epoch 50/100
105/105 [=====] - ETA: 0s - loss: 0.1310 - accuracy: 0.96 - 0s
178us/sample - loss: 0.0674 - accuracy: 0.9810
Epoch 51/100
105/105 [=====] - 0s 206us/sample - loss: 0.0652 - accuracy: 0.9810
Epoch 52/100
105/105 [=====] - 0s 208us/sample - loss: 0.0623 - accuracy: 0.9810
Epoch 53/100
105/105 [=====] - 0s 229us/sample - loss: 0.0633 - accuracy: 0.9810
Epoch 54/100
105/105 [=====] - 0s 179us/sample - loss: 0.0618 - accuracy: 0.9810
Epoch 55/100
105/105 [=====] - 0s 166us/sample - loss: 0.0734 - accuracy: 0.9810
Epoch 56/100
105/105 [=====] - 0s 217us/sample - loss: 0.0979 - accuracy: 0.9429
Epoch 57/100
105/105 [=====] - 0s 169us/sample - loss: 0.0703 - accuracy: 0.9810
Epoch 58/100
105/105 [=====] - 0s 177us/sample - loss: 0.0753 - accuracy: 0.9714
Epoch 59/100
105/105 [=====] - 0s 131us/sample - loss: 0.0942 - accuracy: 0.9524
Epoch 60/100
105/105 [=====] - 0s 246us/sample - loss: 0.0859 - accuracy: 0.9714
Epoch 61/100
105/105 [=====] - 0s 224us/sample - loss: 0.1255 - accuracy: 0.9524
Epoch 62/100
105/105 [=====] - 0s 198us/sample - loss: 0.0627 - accuracy: 0.9714
Epoch 63/100
105/105 [=====] - 0s 228us/sample - loss: 0.0963 - accuracy: 0.9429
Epoch 64/100
105/105 [=====] - 0s 251us/sample - loss: 0.1181 - accuracy: 0.9429
Epoch 65/100
105/105 [=====] - 0s 220us/sample - loss: 0.1049 - accuracy: 0.9524
Epoch 66/100
105/105 [=====] - 0s 296us/sample - loss: 0.0535 - accuracy: 0.9810
Epoch 67/100
105/105 [=====] - 0s 189us/sample - loss: 0.0868 - accuracy: 0.9714
Epoch 68/100
105/105 [=====] - 0s 256us/sample - loss: 0.0919 - accuracy: 0.9429
Epoch 69/100
105/105 [=====] - 0s 217us/sample - loss: 0.0648 - accuracy: 0.9810
Epoch 70/100
105/105 [=====] - 0s 184us/sample - loss: 0.0751 - accuracy: 0.9810
Epoch 71/100
105/105 [=====] - 0s 229us/sample - loss: 0.0775 - accuracy: 0.9714
Epoch 72/100
105/105 [=====] - 0s 301us/sample - loss: 0.0706 - accuracy: 0.9810
Epoch 73/100
105/105 [=====] - 0s 112us/sample - loss: 0.0813 - accuracy: 0.9714
Epoch 74/100
105/105 [=====] - 0s 96us/sample - loss: 0.0643 - accuracy: 0.9810
Epoch 75/100
105/105 [=====] - 0s 91us/sample - loss: 0.0610 - accuracy: 0.9810
Epoch 76/100
105/105 [=====] - 0s 166us/sample - loss: 0.0690 - accuracy: 0.9810
Epoch 77/100
105/105 [=====] - 0s 147us/sample - loss: 0.0800 - accuracy: 0.9810
Epoch 78/100
105/105 [=====] - ETA: 0s - loss: 0.0707 - accuracy: 0.96 - 0s
152us/sample - loss: 0.0770 - accuracy: 0.9714
Epoch 79/100
105/105 [=====] - 0s 129us/sample - loss: 0.0689 - accuracy: 0.9810
```

```

105/105 [=====] - 0s 129us/sample - loss: 0.0669 - accuracy: 0.9810
Epoch 80/100
105/105 [=====] - 0s 128us/sample - loss: 0.0736 - accuracy: 0.9810
Epoch 81/100
105/105 [=====] - 0s 110us/sample - loss: 0.0643 - accuracy: 0.9810
Epoch 82/100
105/105 [=====] - 0s 145us/sample - loss: 0.0616 - accuracy: 0.9810
Epoch 83/100
105/105 [=====] - 0s 123us/sample - loss: 0.0617 - accuracy: 0.9810
Epoch 84/100
105/105 [=====] - 0s 235us/sample - loss: 0.0634 - accuracy: 0.9810
Epoch 85/100
105/105 [=====] - 0s 142us/sample - loss: 0.0629 - accuracy: 0.9810
Epoch 86/100
105/105 [=====] - 0s 150us/sample - loss: 0.0607 - accuracy: 0.9810
Epoch 87/100
105/105 [=====] - 0s 152us/sample - loss: 0.0586 - accuracy: 0.9810
Epoch 88/100
105/105 [=====] - 0s 117us/sample - loss: 0.0674 - accuracy: 0.9810
Epoch 89/100
105/105 [=====] - 0s 128us/sample - loss: 0.0830 - accuracy: 0.9619
Epoch 90/100
105/105 [=====] - 0s 154us/sample - loss: 0.0950 - accuracy: 0.9429
Epoch 91/100
105/105 [=====] - 0s 92us/sample - loss: 0.1034 - accuracy: 0.9619
Epoch 92/100
105/105 [=====] - 0s 154us/sample - loss: 0.1028 - accuracy: 0.9524
Epoch 93/100
105/105 [=====] - 0s 134us/sample - loss: 0.1106 - accuracy: 0.9619
Epoch 94/100
105/105 [=====] - 0s 144us/sample - loss: 0.1662 - accuracy: 0.9429
Epoch 95/100
105/105 [=====] - 0s 152us/sample - loss: 0.1325 - accuracy: 0.9524
Epoch 96/100
105/105 [=====] - 0s 141us/sample - loss: 0.0561 - accuracy: 0.9810
Epoch 97/100
105/105 [=====] - 0s 90us/sample - loss: 0.0664 - accuracy: 0.9810
Epoch 98/100
105/105 [=====] - 0s 111us/sample - loss: 0.0879 - accuracy: 0.9619
Epoch 99/100
105/105 [=====] - 0s 95us/sample - loss: 0.0499 - accuracy: 0.9810
Epoch 100/100
105/105 [=====] - 0s 138us/sample - loss: 0.0760 - accuracy: 0.9619


```

In [43]:

```

plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['loss'], label='loss')
plt.legend()
plt.show()

```



Evaluation of Learning Model

In [44]:

```

y_hat = model.predict(X_test)

```

In [45]:

```
In [45]:
```

```
model.evaluate(X_test, y_test)
```

```
45/45 [=====] - 0s 911us/sample - loss: 0.0628 - accuracy: 0.9556
```

```
Out[45]:
```

```
[0.06283027674588892, 0.95555556]
```

Confusion Matrix

```
In [46]:
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

Confusion matrix C is such that $c(i,j)$ is equal to the number of observations known to be in group i but predicted to be in group j .

```
In [47]:
```

```
y_test_class = np.argmax(y_test,axis=1)
y_pred_class = np.argmax(y_hat,axis=1)
```

```
In [48]:
```

```
print(classification_report(y_test_class,y_pred_class))
C = confusion_matrix(y_test_class,y_pred_class)
print(C)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	0.94	0.94	0.94	16
2	0.93	0.93	0.93	14
avg / total	0.96	0.96	0.96	45


```
[[15  0  0]
 [ 0 15  1]
 [ 0  1 13]]
```

```
In [49]:
```

```
heatmap(C, square=True, annot=True, cbar=False, cmap="YlGnBu")
plt.xlabel('predicted value')
plt.ylabel('true value')
```

```
Out[49]:
```

```
Text(91.68,0.5,'true value')
```

